

CSE 3100: Systems Programming

Lecture 7 : More on Memory, Pointers and Structures

Professor Kaleel Mahmood

Department of Computer Science and Engineering

University of Connecticut

A Few Administrative Announcements:

Don't Cheat.

Paid Summer Research Program in Security at Uconn:
https://khan.engr.uconn.edu/reu_site/index.html



Welcome to the **BEACH**



Belonging, Equity, Affinity,
Community Hangout




ITE Room 360 (follow the signs!)

CSE is committed to building a culture of **belonging and inclusion** for our students. The BEACH is a safe, **open space** dedicated to collaboration for students with the key purpose of **constructing your own community** of support in the department and at UConn.

The BEACH offers the following:

- Dedicated **open space** for students to organically come together during drop-in hours, separate from “traditional” faculty or TA office hours. **Collaborate**, share research and internship advice, or **hang out** over our collection of board games!
- Head TAs, or “**Lifeguards**,” to support students and other CSE TAs.
- Affinity blocks for groups and clubs to **reserve meeting and event space**.
- **Collab Lab hours** for grads and undergrads to mingle while working and introduce UGs to research through grad student seminars and informal discussions.



UG Open Hours

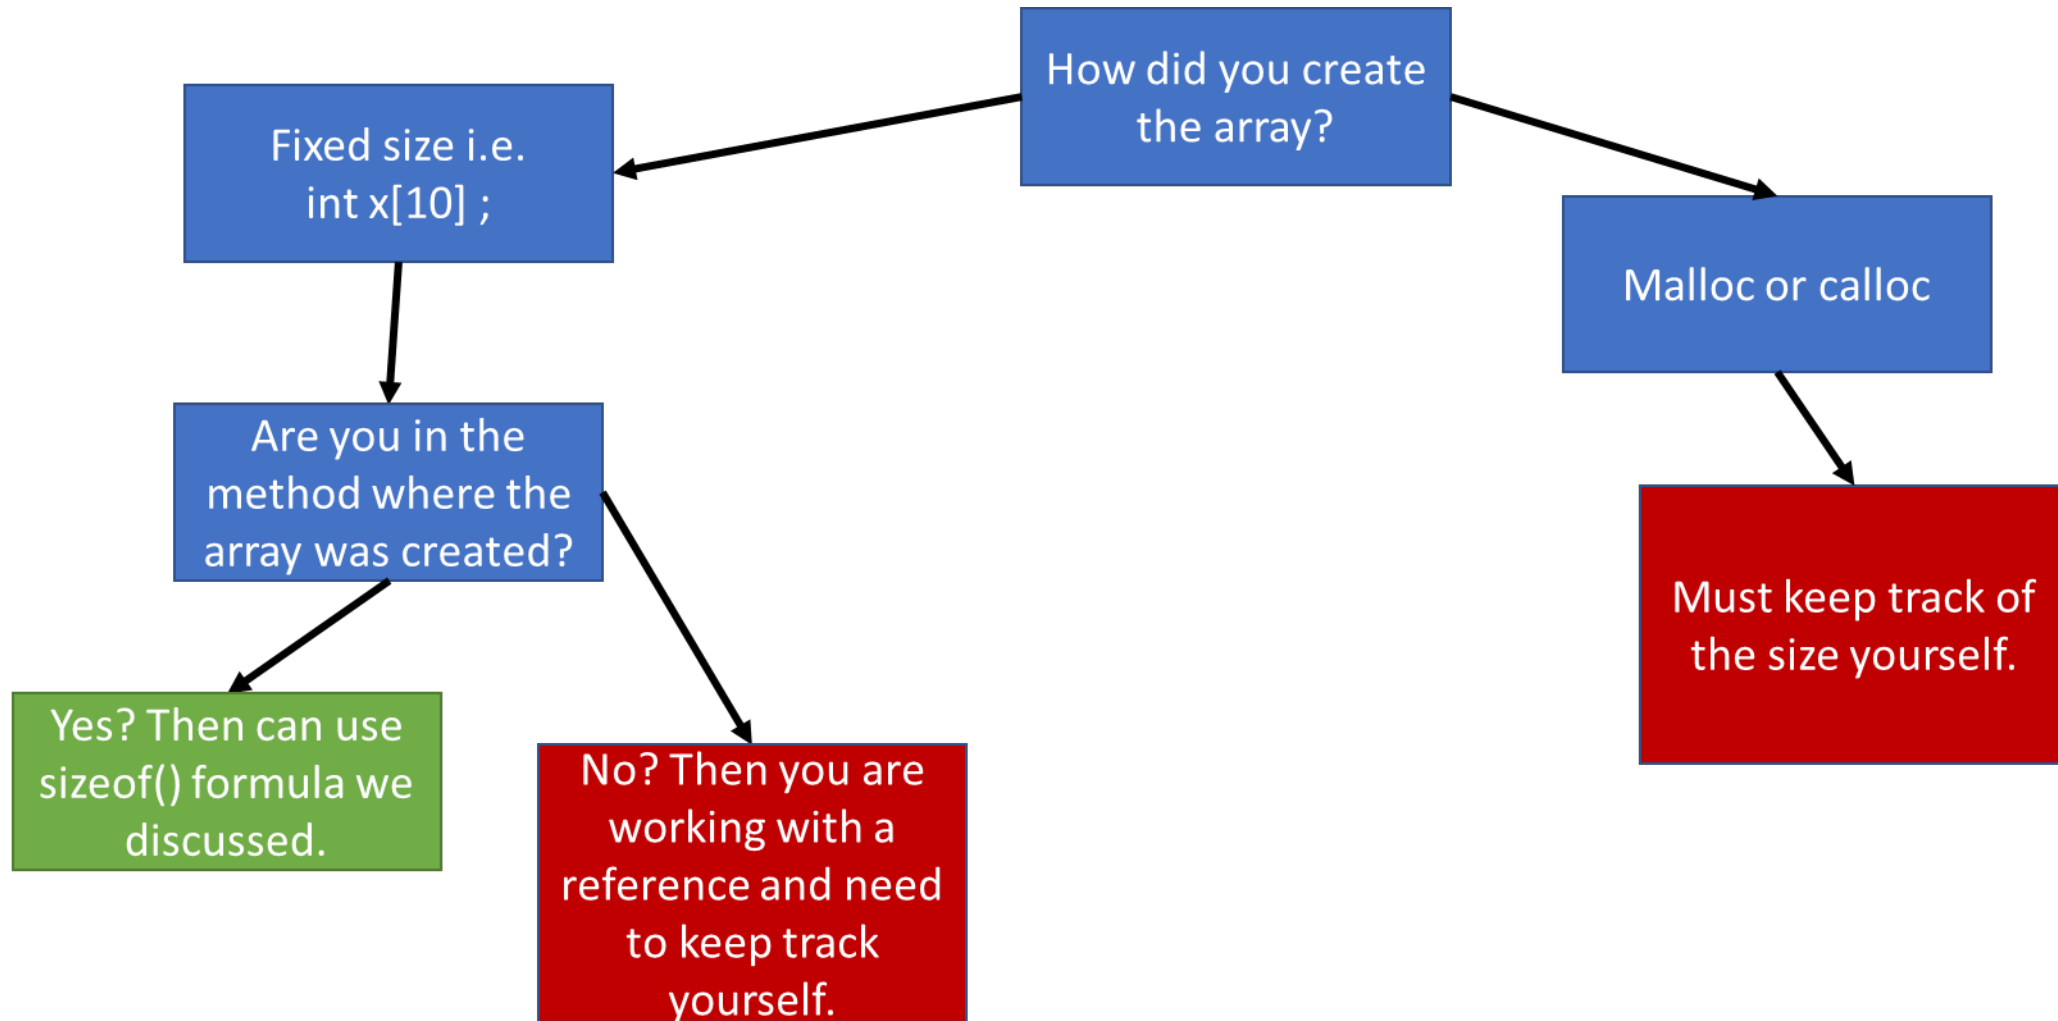
Mon-Fri
4pm-8pm



Computer Science & Engineering

ITE Room 360 (follow the signs!)

Review from last lecture: Memory Allocation



Review from last lecture: Memory Allocation

- Two ways to create arrays:

1. `int x[3];`

2. `int *x = malloc(sizeof(int) * 3)`

In functions why would we want to use malloc over the “[]” approach?

In functions why would we want to use malloc over the “[]” approach?

Car Salesman: *slaps roof of heap*
“This bad boy can hold so many arrays in it”




- When we declare arrays in functions with the “[]” approach that memory is allocated from the stack.
- The stack is a FIXED size, so we could actually run out of memory in a function call if we use an especially large “[]” array.
- The heap is not a fixed size and in general larger.
- Also extremely important: we cannot take non-static stack variables outside of the function call, like we can with malloc...

Question: What will this code print out

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 //try to return a stack variable
4 int* AddThreeToArray(int* x) {
5     int solution[3];
6     for (int i = 0; i < 3; i++) {
7         solution[i] = x[i] + 3;
8     }
9     return solution;
10 }
11
12 int main() {
13     //create an array of size 3
14     int* x = malloc(sizeof(int) * 3);
15     x[0] = 1;
16     x[1] = 2;
17     x[2] = 3;
18     // call AddThreeToArray
19     int* solution = AddThreeToArray(x);
20     //Print solution
21     for (int i = 0; i < 3; i++) {
22         printf("solution[%d]=%d\n", i, solution[i]);
23     }
24 }
```

Ideally we want: the solution to be {4,5,6} right?



```
solution[0]=4
solution[1]=188
solution[2]=-858993460
```

When you try to return stack variables from functions:




Will this code work? If so why?

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int* AddThreeToArray(int* x) {
5      int* solution = malloc(sizeof(int) * 3);
6      for (int i = 0; i < 3; i++) {
7          solution[i] = x[i] + 3;
8      }
9      return solution;
10 }
11
12 int main() {
13     //create an array of size 3
14     int* x = malloc(sizeof(int) * 3);
15     x[0] = 1;
16     x[1] = 2;
17     x[2] = 3;
18     // call AddThreeToArray
19     int* solution = AddThreeToArray(x);
20     //Print solution
21     for (int i = 0; i < 3; i++) {
22         printf("solution[%d]=%d\n", i, solution[i]);
23     }
24 }
```

Answer: Yes because malloc takes memory from the heap NOT the stack. Once allocated by us, the heap memory is only freed by us.

```
solution[0]=4
solution[1]=5
solution[2]=6
```



So should we just never use “[]” arrays that come from the stack?



- The answer is not that simple.
- In this class we treat memory as one big block (with three different sections).
- So it seems like it doesn't matter if we are taking memory from the heap or the stack as both are just pieces from the same block.
- In reality stack memory can be accessed faster than the heap. You'll learn about this more in an operating system class probably.

Pointer Arithmetic


Pointers are addresses



- The **value** of a pointer is a **byte address**
 - Unsigned integer used to number bytes in memory
 - Range is between
 - `0x00000000` and `0xFFFFFFFF` [32-bit]
 - `0x0000000000000000` and `0xFFFFFFFFFFFFFFFF` [64-bit]
- If a pointer is an integer, you can do *arithmetic*...
- To *compute* other addresses

Simple Pointer Example: What value is p pointing to?

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main() {
5      //create an array of size 3
6      int* x = malloc(sizeof(int) * 3);
7      x[0] = 7;
8      x[1] = 8;
9      x[2] = 9;
10     //create another pointer
11     int* p = x + 1;
12     printf("Value at pointer p=%d\n", *p);
13 }
```



Value at pointer p=8

There may be some confusion...

- Recall x is an integer pointer to an address in memory.
- $p = x + 1$
- It seems if we add 1 to it we should get **993**.
- But actually p is pointing to **996**.



Address	Value
1020	
1016	
1012	
1008	
1004	
1000	9
996	8
992	7

p points to address 996

x points to address 1000

Adding a pointer and an integer

- Add an integer to a pointer, the result is a pointer of **the same type**
 - It is different from regular integer addition
 - The integer is **automatically scaled** by the size of the type pointed to

- Suppose p is a pointer to type T , and k is an integer


Then both $p + k$ and $k + p$

- Are valid expressions that evaluate to a pointer to type T
- Have a byte address equal to

$(\text{unsigned long})(\text{address stored in } p) + k * \text{sizeof}(T)$

Can we actually see the addresses?

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main() {
5      //create an array of size 3
6      int* x = malloc(sizeof(int) * 3);
7      x[0] = 777;
8      x[1] = 888;
9      x[2] = 999;
10     //create another pointer
11     int* p = x + 1;
12     printf("Address x points to %x\n", x);
13     printf("Address p points to %x\n", p);
14 }
```



Address x points to 41ee0800
Address p points to 41ee0804

Pointer Addition

Suppose p is a pointer to an `int`, and its value is 1000

$p + 1$ is not the next byte address.

It is the address of next item (of type `int`)


Address	Value			To access values
1020		←	$p + 5$	$*(p+5)$ OR $p[5]$
1016		←	$p + 4$	$*(p+4)$ OR $p[4]$
1012		←	$p + 3$	$*(p+3)$ OR $p[3]$
1008		←	$p + 2$	$*(p+2)$ OR $p[2]$
1004		←	$p + 1$	$*(p+1)$ OR $p[1]$
1000		←	$p = 1000$	$*p$ OR $p[0]$
996		←	$p - 1$	$*(p-1)$ OR $p[-1]$
992		←	$p - 2$	$*(p-2)$ OR $p[-2]$

Pointers Subtraction

- Subtract one pointer from another: both must have the **same** type
- The result is **the number of data items** between the two pointers!
 - Not the number of bytes

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int *p = malloc(sizeof(int)*10);
    int *last = p + 9;
    int dist = last - p; // both are
int *
    printf("Distance is %d\n",dist);
    free(p);
    return 0;
}
```



Output

```
$ gcc
ptrsub.c
$ ./a.out
Distance is 9
$
```

Pointer comparisons

- You can also compare two pointers

< > <= >= != ==

- Purpose
 - Check boundary conditions in arrays
 - Manually manage memory blocks
- Semantics
 - Simply based on memory layout!
 - Compare bits in pointers as unsigned integers!

Effect of casting types


- If you cast a pointer type...
 - Any subsequent pointer arithmetic will use the type you chose
- Do not want scaling? Casting a pointer to (**char ***)
 - Because sizeof(char) is 1

```
int * t;  
char * p = (char *) t + 8;           // 8 is not  
scaled  
char * q = (char *) (t + 8); // 8 is scaled
```

Arrays and Pointers 1

- Arrays and pointers can often be used interchangeably

```
int  a[10];  
int *p = a;  
  
// all of the following evaluate to the value of a[0]  
*p           p[0]           *a           a[0]  
  
// all of the following evaluate to the value of a[1]  
*(p+1)       p[1]           *(a+1)       a[1]  
  
// all of the following evaluate to the address of a[0]  
// type is int *  
p            &p[0]          a            &a[0]
```



"array of int" becomes "pointer to int" (array decay)

Arrays and Pointers 2

- Equivalent ways of initializing an array

```
int a[10], *p = a; // not *p = a; it is int *p; p = a;


for(int i=0; i<10; i++) a[i] = i; // array indexing

for(int i=0; i<10; i++) p[i] = i; // indexing via pointer

for(int i=0; i<10; i++) *(p+i) = i; // explicit pointer arithmetic

for(i=0; i<10; i++) i[p] = i; // obfuscated but valid C!

for(i=0; i<10; i++) *p++ = i; // common pointer use idiom
```



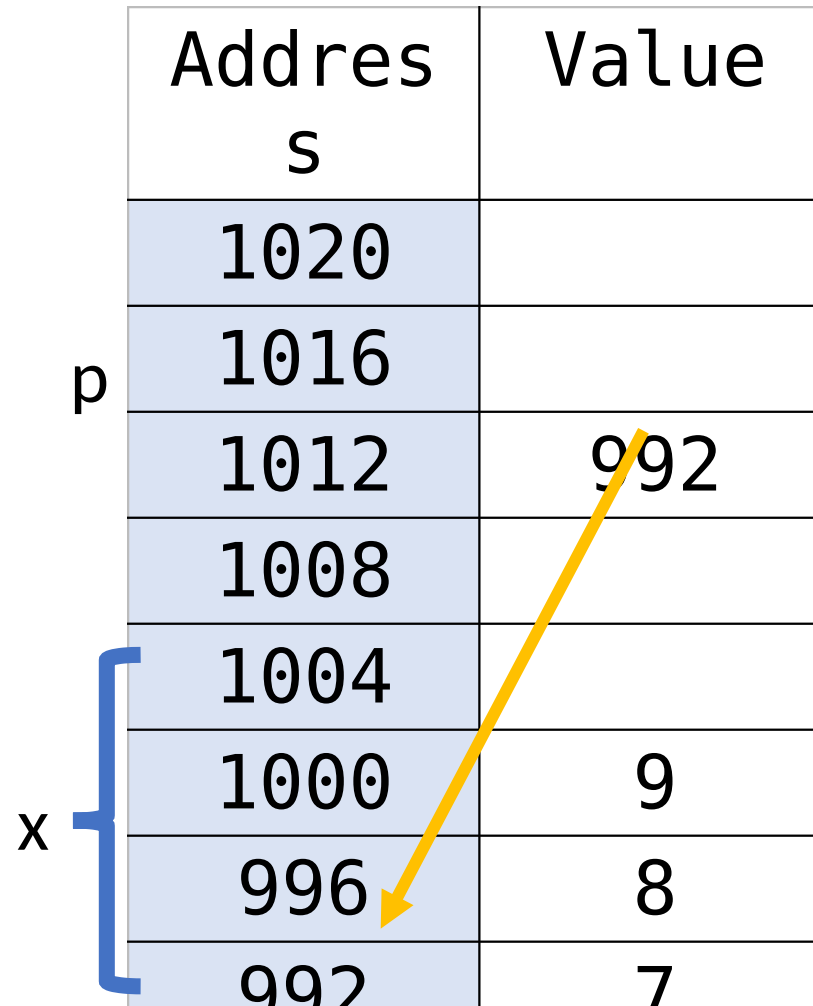
```
int * tp = p;
p ++;
*tp = i;
```

Are arrays and pointers the same?



Visualization of Pointer vs Array

- In this visualization x is an array and p is a pointer.



The diagram illustrates memory addresses and values. A table shows addresses from 1020 down to 992 in descending order. The values at these addresses are 992, 9, 8, and 7 respectively. A pointer 'p' is associated with address 1012, and an array 'x' is associated with addresses 1004, 1000, 996, and 992. A yellow arrow points from the value 992 at address 1012 to the value 8 at address 996.

	Address	Value
	1020	
	1016	
p	1012	992
	1008	
	1004	
	1000	9
x	996	8
	992	7

Arrays and pointers are **NOT** the same

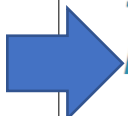
```
int    a[10];
int    *p = a; // a is converted to int *

// a is still an array after &
&a    // pointer to array of 10 int's  int (*)(10)
&p    // pointer to a pointer to int  int **

// a is still an array after sizeof
sizeof(a) // 40 because a is an array of 10 int's
sizeof(p) // 8  because p is a pointer

p++;    // can increment p
a++;    // cannot increment a; this will not compile
        // Similar to n++ vs 2++
```


Code example to further illustrate the point...

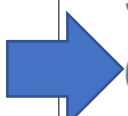


```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main() {
5      //create array
6      int a[] = { 1, 2, 3 };
7      //create pointer to array
8      int* p = a;
9  }
```

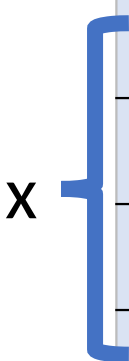
Address	Value
1020	
1016	
1012	
1008	
1004	
1000	
996	
992	

Code example to further illustrate the point...

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main() {
5      //create array
6      int a[] = { 1, 2, 3 };
7      //create pointer to array
8      int* p = a;
9  }
```

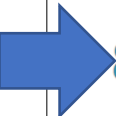


	Address	Value
	1020	
	1016	
p	1012	
	1008	
	1004	
	1000	3
x	996	2
	992	1

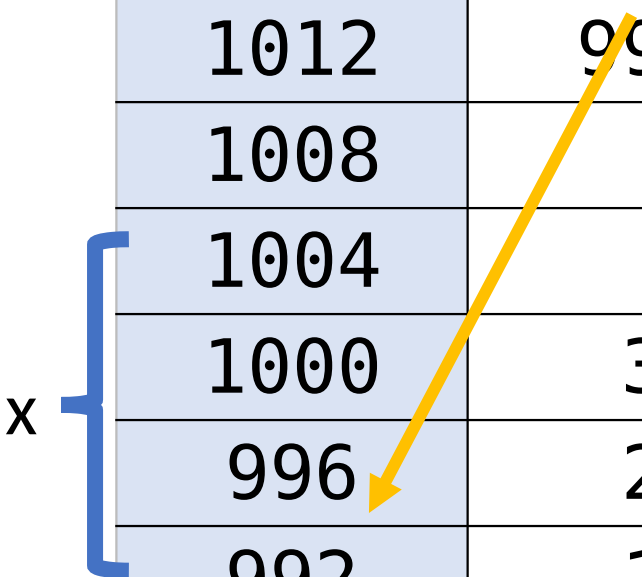


Code example to further illustrate the point...

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main() {
5      //create array
6      int a[] = { 1, 2, 3 };
7      //create pointer to array
8      int* p = a;
9  }
```



	Address	Value
	1020	
	1016	
p	1012	992
	1008	
	1004	
	1000	3
x	996	2
	992	1



Showing the Address for Further Proof


- We can actually print out the address to see that the pointer is pointing to the address AND the array value is stored at that address:

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main() {
5      //create array
6      int a[] = { 1, 2, 3 };
7      //create pointer to array
8      int* p = a;
9      printf("Address of a[0]:%x\n", &a[0]);
10     printf("Address stored in pointer p:%x\n", p);
11 }
```



```
Address of a[0]:e9b6faa8
Address stored in pointer p:e9b6faa8
```

Structures



Structures

```
struct student_grade {  
    char *name;  
    int id;  
    char grade[3];  
};  
...  
struct student_grade student1;  
struct student_grade student2,  
student3;  
struct student_grade all[200];
```

- Mechanism to define new types
 - Also known as “compound types”
 - Used to aggregate related variables of different types
- Structures type declaration
 - Structures can have a **type name**
 - Can have “members” of any type
 - Basic types
 - Pointers
 - Arrays
 - Other structures
- Structure variable definition
 - Specifies **variable name**

Structure Example

```
struct Person {  
    int    age;  
    char  gender;  
};
```

Structure *type declaration*

```
int main(){  
    struct Person p;
```

Structure *variable definition*

```
    p.age = 44;  
    p.gender = 'M';
```

Syntax for field access
similar to Java and Python

```
    struct Person q = {44, 'M'};
```

Structure *variable definition*
and *initialization*

```
    return 0;
```

```
}
```

Example: Array of Structures

```
#include <stdlib.h>

struct Person {
    char    name[32];
    int     age;
    char    gender;
};

int main()
{
    struct Person family[4] = {
        {"Alice", 34, 'F'},
        {"Bob", 40, 'M'},
        {"Charles", 15, 'M'},
        {"David", 13, 'M'}
    };
    int juniorAge = family[3].age;
    return 0;
}
```

- Member name is a char array
- Caveats
 - Names cannot be more than 31 characters long
 - Four persons in family
 - Indexed 0..3
- Array of structures for the whole family
 - Nested initializers

typedef

```
struct Person {  
    char    name[32];  
    int     age;  
    char    gender;  
};  
typedef struct Person TPerson;  
int main()  
{  
    TPerson family[4];  
    ...  
    return 0;  
}
```

```
typedef struct Person {  
    char    name[32];  
    int     age;  
    char    gender;  
} TPerson;
```

- Struct names can become long
- C provide the ability to define type abbreviations
 - **typedef** declaration
 - Give existing type new type name
- Make code more readable
- Structure and typedef declarations often combined

Operations on struct

- Assignment
 - All struct members copied
- Can be passed to functions
 - **By value**
 - Even if some members are arrays!
- Can be returned from a function
- If pass by value, cannot change members in functions
- Passing or returning large structures can be costly

Use pointers to structures!

Structure Alignment

- Structure members are *aligned* for the natural types

Alignment requirements on x64 architecture	
char	1
short	2
int	4
long	8
float	4
double	8

```
struct struct_random {  
    char    x[5]; // bytes 0-4  
    int     y;    // bytes 8-11  
    double  z;    // bytes 16-23  
    char    c;    // byte 24  
};           // Total size 32
```

```
struct struct_sorted {  
    double  z;    // bytes 0 - 7  
    int     y;    // bytes 8 - 11  
    char    x[5]; // bytes 12 - 16  
    char    c;    // byte 17  
};           // Total size 24
```

Lecture Conclusions

Python devs when people start talking about pointers and memory allocation



- As a language C offers an extremely large amount of control over the code you program.
- This also makes it very challenging to learn and pick up.
- You don't want to be like the meme shown on this slide.
- You should understand how to use pointers, arrays and which variables use which memory.
- Understand how to create structures in C.

Figure Sources

1. <https://i.kym-cdn.com/photos/images/newsfeed/001/608/838/9f2.jpg>
2. <https://img.memegenerator.net/instances/62930796.jpg>
3. https://i.kym-cdn.com/entries/icons/original/000/028/596/dsmGaKWMcHxe9QuJtq_ys30PNfTGnMsRuHuo_MUzGCg.jpg
4. https://mediaproxy.salon.com/width/1200/https://media.salon.com/2013/07/2001_monolith.jpg
5. <https://programmerhumor.io/wp-content/uploads/2023/01/programmerhumor-io-python-memes-backend-memes-e1c8fbd8d3296f0-608x596.jpg>